

Second Order Greeks via Adjoint Algorithmic Differentiation

UIUC MSFE

University of Illinois Urbana-Champaign

Junru Wang Zhaolin Shi Zhengtao Yu
Yaru Guan Zilin Mao Wenqi Guo

Dec 3, 2025

J.P.Morgan



Corporate Advisor: Gan Wang, Vicky Luo

Sponsor Company: J.P. Morgan

Instructor: Prof. Feng

Project Motivation:

- **Problem:** Traditional finite-difference bumping scales poorly—computing n Greeks requires $O(n)$ re-pricings.
- **Solution:** Reverse-mode AAD computes all first-order Greeks in one backward sweep; Hessian AAD (Edge-Pushing & Taylor Expansion) extends this to second-order Greeks.
- **Result:** Achieved great speedup over bumping on both PDE and Monte Carlo pricing engines.

Insight of “Automatic”

- Programs = compositions of **elementary operations**
- Derivatives of each operation are known
- AD applies the **chain rule** systematically

Modes of Automatic Differentiation

• Forward Mode

- Propagates derivatives alongside values
- Efficient when number of inputs n is small
- Computes one directional derivative per pass

• Reverse Mode

- Propagates adjoints backward from output
- Superior when $\text{inputs} \gg \text{outputs}$
- Ideal for financial models: option pricing, risk, Greeks




Figure 1: Expression graph for the log-normal density

$$f(y; \mu, \sigma) = \frac{1}{y\sigma\sqrt{2\pi}} \exp\left(-\frac{(\ln y - \mu)^2}{2\sigma^2}\right)$$

Computation Graph

Edge

- An edge represents a function argument (and also a data dependency).
- Edges are simply pointers to nodes.

Node


- The head node of an edge is a function of that edge's tail node.
- A node “knows” how to compute its value and the value of its derivative w.r.t. each argument (edge) multiplied by a derivative arriving from an arbitrary input, i.e.,

$$\frac{\partial F}{\partial u} = \frac{\partial F}{\partial v} \cdot \frac{\partial v}{\partial u}.$$

Tape

- ❶ Topological indices of nodes (predecessor list)
- ❷ Forward values
- ❸ Numerical values of Jacobians / gradients

Computation graphs are usually a directed acyclic graph (DAG).



$$f(y; \mu, \sigma) = \frac{1}{y\sigma\sqrt{2\pi}} \exp\left(-\frac{(\ln y - \mu)^2}{2\sigma^2}\right)$$

First-order Greeks

- Finite-difference bumping computes each Greek separately:

$$\Gamma_i \approx \frac{F(x + he_i) - F(x)}{h},$$

- Reverse-mode AAD yields *all* first-order Greeks in a single backward sweep:

Time complexity

- Finite-difference bumping: $\mathcal{O}(n)$ evaluations of the pricing function.
- Reverse-mode AAD: $\mathcal{O}(1)$ backward sweep (independent of n).
- Thus, when the number of inputs is large, AAD decisively outperforms bumping.

Beyond first order: towards Hessians

- To extend AAD-style efficiency to second-order Greeks (Hessians), we draw on the framework of Gower & Mello (2012), *A new framework for the computation of Hessians*.
- This introduces the **edge-pushing algorithm**, which implements the chain rule for second-order derivatives and propagates Hessian through the graph.
- Our benchmark method for comparison remains finite-difference bumping.

Function f and its Hessian f''

$$f(x) = e_\ell^\top (\Phi_\ell \circ \dots \circ \Phi_1)(P^\top x)$$

where

- $x \in \mathbb{R}^n$: input vector
- $P = [I_n \ 0]$: embedding matrix, $P^\top x = (x, 0, \dots, 0) \in \mathbb{R}^{n+\ell}$
- Φ_i : i -th state transformation, updates only one register (one component) of the state vector
- e_ℓ^\top : selects the final register (output)

$$f''(x) = P \left(\sum_{i=1}^{\ell} \left(\prod_{j=1}^{i-1} (\Phi'_j)^\top \right) ((\bar{v}^i)^\top \Phi''_i) \left(\prod_{j=1}^{i-1} \Phi'_{i-j} \right) \right) P^\top$$

where

- Φ'_j : Jacobian of Φ_j , Φ''_j : Hessian
- \bar{v}^i : adjoint vector (backpropagated weight)
- P, P^\top : map between input space and extended state space

Hessian's Chain Rule

$$f'' = P \sum_{i=1}^{\ell} \left(\prod_{j=1}^{i-1} (\Phi'_j)^\top \right) \left((\bar{v}^i)^\top \Phi''_i \right) \left(\prod_{j=1}^{i-1} \Phi'_{i-j} \right) P^\top$$

Block form:

$$f'' = PWP^\top, \quad \text{with } W = \sum_{i=1}^{\ell} W_i$$

Previous summand W_{i-1} :

$$W_{i-1} = \underbrace{\left((\Phi'_1)^\top \cdots (\Phi'_{i-2})^\top \right)}_{\text{left factors}} \left((\bar{v}^{i-1})^\top \right) \underbrace{\left(\Phi''_{i-1} \right)}_{\text{local Hessian at } i-1} \underbrace{\left(\Phi'_{i-2} \cdots \Phi'_1 \right)}_{\text{right factors}}$$

Current summand W_i :

$$W_i = \underbrace{\left((\Phi'_1)^\top \cdots (\Phi'_{i-2})^\top \right)}_{\text{same as before}} \underbrace{\left(\Phi'_{i-1} \right)^\top}_{\text{extra}} \left((\bar{v}^i)^\top \right) \underbrace{\left(\Phi''_i \right)}_{\text{local Hessian at } i} \underbrace{\Phi'_{i-1}}_{\text{extra}} \underbrace{\left(\Phi'_{i-2} \cdots \Phi'_1 \right)}_{\text{same as before}}$$

Core idea: Each node performs the sequence **Push** → **Create** → **Update** to incrementally construct the Hessian, with each sweep adding one more W_i summand to W .

1. Push Step — This allows W to gain one additional summand.

$$W \leftarrow (\phi'_i)^T W \phi'_i$$

Propagates existing second-order information to the parent nodes through the local Jacobian.

2. Create Step

$$W \leftarrow W + \bar{v}^T \phi''_i$$

Introduces new second-order contributions at the current node, arising from its local Hessian.

3. Update Step

$$\bar{v}^T \leftarrow \bar{v}^T \phi'_i$$

Performs the standard first-order adjoint update (reverse-mode propagation).

Componentwise Form Applied to a Computation Graph

- **Sweeping node 3:** A new nonlinear arc $\{1, 2\}$ is created.
- **Sweeping node 2:** The arc $\{1, 2\}$ is pushed and split into
 - $\{0, 1\}$ with weight $1 \cdot 2v_0$,
 - $\{-1, 1\}$ with weight $1 \cdot 3$.
- **Sweeping node 1:** The arc $\{-1, 1\}$ is pushed and split into
 - $\{-2, -1\}$,
 - $\{-1, -1\}$ with weight $2 \cdot 3 \cdot e^{v-1}$.




Figure. Edge-pushing applied to a computational graph of $f(x) = (x_{-2} + e^{x-1})(3x_{-1} + x_0^2)$.

Taylor Expansion as an Alternative Hessian Engine

1. Taylor Forward (local expansion)

- Each operator is expanded to second order:

$$f(z + \varepsilon) \approx f(z) + f'(z)\varepsilon + \frac{1}{2}f''(z)\varepsilon^2.$$

- Forward sweep propagates (v_0, v_1, v_2) at every node.
- All second-order monomials $(\varepsilon_i^2, \varepsilon_i\varepsilon_j)$ generated locally.

2. Taylor Backpropagation (one reverse sweep)

- Forward sweep already provides derivatives of intermediates.
- One reverse sweep folds all monomials back to inputs:

$$H_{ij} = \sum_{k,m} \frac{\partial^2 V}{\partial u_k \partial u_m} \frac{\partial u_k}{\partial x_i} \frac{\partial u_m}{\partial x_j} + \dots$$

3. Comparing Hessian Engines: Edge-Pushing vs Taylor

Aspect	Edge-Pushing (EP)	Taylor Expansion
Derivative storage	second-order edge pairs	local (v_0, v_1, v_2) per node
Reverse sweep cost	requires large nonlinear graph	single sweep
Scaling with steps	graph grows with time	linear in depth
Hessian symmetry	needs explicit enforcement	automatic

Takeaway: Taylor avoids graph blow-up, preserves symmetry, and provides scalable second-order derivatives.

Implementation of the Taylor Expansion Engine

1. Taylor AD Engine

- Each `ADVar` node stores (v_0, v_1, v_2) .
- Overloaded ops (+, -, *, exp, log, sqrt, norm_cdf, ...) propagate local second-order Taylor terms through the graph.
- `taylor_backpropagate(target)`: a single reverse sweep returns the full Hessian w.r.t. $(S_0, \sigma, K, r, T, \dots)$.

2. BSM 5D Benchmark (accuracy & speed)

Black–Scholes 5D setting with inputs (S_0, σ, K, r, T) ; Edge–Pushing (EP) used as the reference Hessian.

Hessian accuracy vs EP

Method	Frobenius rel. error	Comment
Taylor	1.82e–02	matches EP structure
Finite Difference (FD)	1.83e–02	slightly noisier

Runtime (baseline = EP)

Method	Time (ms)	vs EP
Edge–Pushing (EP)	30.998	1.00×
Taylor	292.687	0.11×
Finite Difference (FD)	0.105	295.92×

Taylor achieves EP-level accuracy, with a predictable single-tape structure. FD is extremely fast but less robust for higher-dimensional Greeks.

1. Embedding Taylor Expansion into the PDE Solver

- Treat PDE coefficients a, b, c as differentiable functions of AD variables.
- Each grid state u^n becomes an **ADVar** that carries (v_0, v_1, v_2) .
- Crank–Nicolson update remains inside the AD graph:

$$u^{n+1} = A^{-1}(Bu^n + b).$$

- One call to `taylor.backpropagate()` returns the entire gradient and Hessian w.r.t. any model parameter (no additional adjoint rules required).

2. Extension to Local Volatility (Dupire SVI)

- Replace constant σ with Dupire local vol $\sigma_{\text{loc}}(K, T)$ from an SVI surface.
- PDE coefficients become nonlinear in (S, t) and depend on SVI parameters $\{a_i, b_i, \rho_i, m_i, \eta_i\}$.
- Taylor automatically differentiates all PDE coefficients, preserving Hessian symmetry under highly nonlinear local vol dynamics.

Takeaway: Taylor provides a fully embedded second–order engine inside the PDE solver, enabling gradients and Hessians w.r.t. all SVI parameters in a single pass.

PDE + Taylor + SVI Results (BSM Setting)

Dupire–SVI local volatility. Grid $M = 121$, $N_{\text{base}} = 100$. 25 SVI parameters $\{a_i, b_i, \rho_i, m_i, \eta_i\}_{i=0}^4$.

1. Gradient benchmark: Taylor vs FD (25 parameters)

Metric	Taylor (AD)	Bumping FD	Note
Avg relative error	2.6×10^{-3}	—	high alignment
Max relative error	1.65×10^{-1} (on a_1)	—	localized worst case
One-shot AD time	19.8 s	—	25 greeks at once
FD total (25 bumps)	—	8.79 s	1 greek per run
Speedup (FD/AD)	0.44×		FD faster (gradients)

2. Full 25×25 SVI Hessian: Taylor vs FD

Metric	Taylor (AD)	FD Bumping	Note
Avg relative error	1.26×10^{-2}	—	stable across entries
Max relative error	8.65×10^{-1}	—	one SVI axis
One-shot AD Hessian	20.65 s	—	dense 25×25
FD total (25×25 bumps)	—	141.47 s	625 valuations
Speedup (FD/AD)	6.85× faster		AD faster (Hessian)

Conclusion: Taylor more costly for gradients, but for **full Hessian** it is $\sim 7\times$ faster than bumping FD with strong accuracy.

Stage	Implementation	Speedup	Status
Stage 0	Pure Python	1.0× (baseline)	Too slow
Stage 1	Cython + Python dict	12.0×	✓ Deployed
Stage 2	C++ unordered_map	3.2×	Alternative

Key Technologies:

- **Cython:** Hybrid Python/C extension that compiles Python-like code to C, providing near-C performance while maintaining Python's ease of use. Enables type declarations and direct C API calls.
- **Python dict:** Native Python dictionary with $O(1)$ average-case hash lookups. Leverages highly optimized CPython implementation for edge storage in sparse Hessian graphs.
- **C++ unordered_map:** STL hash table with comparable performance but higher implementation overhead. Pure C++ alternative when full compilation is needed.

Key Finding: Cython + Python dict achieves **12× speedup** with minimal complexity, outperforming pure C++ alternatives.

B-spline Calibration on SPX Options Data

Real Market Data: SPX Options

- **Dataset:** SPY options expiring 2025-02-06
- **Underlying:** S&P 500 Index, $S_0 = \$6081.76$
- **Options:** 5,405 valid contracts
- **Maturities:** Multiple expiration dates

B-spline Parametrization


The calibrated B-spline weights

$\mathbf{w} = (w_1, \dots, w_m)$ define a smooth local volatility surface:

$$\sigma_{\text{loc}}(K, T; \mathbf{w}) = \sum_{i=1}^m w_i B_i(K, T),$$

where $B_i(K, T)$ are tensor-product cubic B-spline basis functions.


Volatility Smiles: Market vs B-spline Fit



Calibration Method:

- **Optimizer:** L-BFGS-B with AAD gradients
- **Objective:** Minimize pricing error between market and model prices
- **EP Hessian:** Used for validation and second-order optimization

Calibrated B-spline Volatility Surface



Surface Properties:

- **Smoothness:** C^2 continuity across the entire (K, T) domain
- **Locality:** Each basis function has compact support \Rightarrow sparse Hessian
- **Linearity:** σ_{loc} depends linearly on weights $\mathbf{w} \Rightarrow$ PDE coefficients have simple \mathbf{w} -dependence
- **Efficiency:** Gradient and Hessian w.r.t. \mathbf{w} computed via AAD Edge-Pushing

1. B-spline parametrization of the model

-

$$\sigma_{\text{loc}}(t, x; \mathbf{w}) = \sum_{i=1}^m w_i B_i(x),$$

where $\mathbf{w} = (w_1, \dots, w_m)$ are the spline weights (calibration parameters).

- This gives a smooth, low-dimensional and *linear* dependence of the PDE coefficients on \mathbf{w} .

2. PDE structure and dependence on B-splines

$$\frac{\partial V}{\partial t} + \frac{1}{2} \sigma_{\text{loc}}^2(t, x; \mathbf{w}) x^2 \frac{\partial^2 V}{\partial x^2} + rx \frac{\partial V}{\partial x} - rV = 0,$$
$$A(\mathbf{w}) V^{n+1} = B(\mathbf{w}) V^n,$$

where all matrix entries depend on \mathbf{w} *only through the spline expansion*.

- The dependence is smooth and local: at each grid point x_j ,

$$\sigma_{\text{loc}}(t_n, x_j; \mathbf{w}) = \sum_{i: B_i(x_j) \neq 0} w_i B_i(x_j),$$

$$\mathbf{w} \longmapsto V(0, x_0; \mathbf{w}),$$

Numerical Verification: EP vs Taylor vs Bumping2

Results: 9-Parameter System ($n_{\text{knots}} = 5$)

Metric	Edge-Pushing	Taylor	Bumping2	Best Method
Price	13.0193	13.0193	13.0193	All match
Max Element Error	—	—	—	<0.1%
PDE Solved	1	1	180	EP/Taylor: 180× fewer
Computation Time	93.6 sec	412.3 sec	1252 sec	EP: 13.4× faster


Results: 14-Parameter System ($n_{\text{knots}} = 10$)

Metric	Edge-Pushing	Taylor	Bumping2	Best Method
Price	12.0485	12.0485	12.0485	All match
Max Element Error	—	—	—	<0.35%
PDE Solved	1	1	420	EP/Taylor: 420× fewer
Computation Time	148 sec	885 sec	4,701 sec	EP: 32× faster

Results: 19-Parameter System ($n_{\text{knots}} = 15$)

Metric	Edge-Pushing	Taylor	Bumping2	Best Method
Price	12.3402	12.3402	12.3402	All match
Max Element Error	—	—	—	<0.35%
PDE Solved	1	1	760	EP/Taylor: 760× fewer
Computation Time	173 sec	1221 sec	6451 sec	EP: 42× faster

Edge-Pushing Hessian Visualization



10×10 Hessian for B-spline PDE Calibration

- **Sparse structure:** Only 36% non-zero elements due to B-spline compact support
- **Edge-Pushing:** Exploits sparsity to compute full Hessian in one PDE solve
- **Validation:** EP vs numerical Hessian error $< 10^{-8}$ (relative Frobenius norm)
- **Speedup:** 42× faster than Bumping2 for 19-parameter system

Pathwise Differentiation for Basket Option

1. Model Setup

$$B_T = \sum_{i=1}^n w_i S_{T,i}, \quad P = e^{-rT} (B_T - K)^+, \quad V = \mathbb{E}[P(\mathbf{S}_T(\theta))].$$

2. Pathwise Estimators (The Chain Rule) Interchanging differentiation and expectation:

$$\text{Gradient: } \frac{\partial V}{\partial \theta_k} = \mathbb{E} \left[\frac{\partial P}{\partial B_T} \frac{\partial B_T}{\partial \theta_k} \right]$$

$$\text{Hessian: } \frac{\partial^2 V}{\partial \theta_k \partial \theta_m} = \mathbb{E} \left[\frac{\partial^2 P}{\partial B_T^2} \frac{\partial B_T}{\partial \theta_k} \frac{\partial B_T}{\partial \theta_m} + \frac{\partial P}{\partial B_T} \frac{\partial^2 B_T}{\partial \theta_k \partial \theta_m} \right]$$

3. Payoff Derivatives For the Call Option Payoff:

- **1st Order:** $\frac{\partial P}{\partial B_T} = e^{-rT} \mathbf{1}_{\{B_T > K\}}$ (Digital Payoff).

- **2nd Order:** $\frac{\partial^2 P}{\partial B_T^2} = e^{-rT} \delta(B_T - K).$

(Note: Handling the Dirac mass δ requires smoothing.)

Softplus Smoothing


Smooth surrogate for the hard payoff:

$$\text{softplus}_\alpha(x) = \frac{1}{\alpha} \ln(1 + e^{\alpha x}) \approx \max(0, x) \quad (\alpha \uparrow)$$

- $\alpha \uparrow$: **sharper** kink, closer to hard payoff; price bias \downarrow but curvature near $x \approx 0$ increases \Rightarrow MC variance / numerical fragility \uparrow .
- $\alpha \downarrow$: **smoother** kink, further from hard payoff; price bias \uparrow but typically lower variance and more stable second derivatives.

Softplus bias (vs. ReLU): The maximum pointwise error occurs at the money ($x = 0$):

$$b_\alpha^{\max} = \text{softplus}_\alpha(0) - \max(0, 0) = \frac{\ln 2}{\alpha} \approx \frac{0.693}{\alpha}.$$



Note: domain clipped to $[-0.25, 0.25]$.

Representative α results:

α	Γ rel	Γ mean $ \Delta $	V_n rel	V_n mean $ \Delta $	V_l rel	V_l mean $ \Delta $
10.00	8.94e-03	1.41e-06	6.62e-05	7.78e-07	2.68e-05	1.14e-05
30.00	4.86e-03	1.02e-06	1.17e-04	1.25e-06	3.40e-04	1.31e-04
100.00	5.40e-03	1.05e-06	1.60e-03	1.64e-05	4.48e-03	1.71e-03

Verdict: Choose $\alpha = 30$.

- Best trade-off: accuracy, variance and speed.
- EP vs. FD speedup: **1.13 \times** .


Key stats at $\alpha = 30$:

- Γ rel err: 4.86×10^{-3}
- Vanna/Volga rel err: $\sim 10^{-4}$

Context & Trade-offs:

- **Small α (e.g., 10):**
Higher Vanna/Volga stability, but higher bias and slower execution (0.94 \times).
- **Large α (e.g., 100):**
Fastest speed (1.18 \times), lower bias but errors increase significantly on account of variance.

EP computation graph for Basket option



- H_{all} is the $2n \times 2n$ Hessian with respect to $[S_vars; Sig_vars]$ on a single tape.
 - Gamma, Vanna and Volga come from the three blocks:

$$\text{Gamma}(S, S) = H_{\text{all}}[0:n, 0:n],$$

$$\text{Vanna}(S, \sigma) = H_{\text{all}}[0:n, n:2n],$$

$$\text{Volga}(\sigma, \sigma) = H_{\text{all}}[n:2n, n:2n].$$

EP vs FD: Accuracy and Speed on Basket option

Accuracy at $n = 4$ (alpha=30, 10000 paths, 50 steps)

Greeks block	Frobenius rel. error	mean $ \Delta $
Gamma (S, S)	4.52e-03	1.54e-06
Vanna (S, σ)	1.10e-04	3.40e-06
Volga (σ, σ)	2.43e-04	4.06e-04

EP Greeks track the FD baseline very closely; discrepancies are at the level of Monte Carlo noise.

Timing vs dimension (alpha=30, 10000 paths, 50 steps).

n assets	EP time (s)	FD time (s)	Speedup (FD / EP)
5	35	20	0.59×
7	46	49	1.07×
10	64	131	2.07×
15	95	412	4.36×
20	133	924	6.95×
25	169	1757	10.42×
30	206	2864	13.93×
40	298	6517	21.88×

As dimension n grows, FD cost scales combinatorially in the number of bumps, while EP keeps a single-tape Hessian and benefits from relatively better scaling.

Conditional Expectation & Vibrato

CE idea: Move the non-smoothness away in the direction of a chosen asset. Use the last-step conditional expectation conditional on $S_{N-1}^{(k)} = x$:

$$\begin{aligned}\phi_k(x) &= \mathbb{E}[f(S_T) \mid S_{N-1}^{(k)} = x] \\ &= \int_{\mathbb{R}^d} f(z) p_T(z \mid x) dz.\end{aligned}$$

- $p_T(z \mid x)$: conditional joint density of $S_T = z$ given the (k th) asset satisfies $S_{N-1}^{(k)} = x$.
- f may be non-smooth in the basket direction, but $p_T(\cdot \mid x)$ is C^∞ in x and differentiable.
- $\phi_k(x)$ can be viewed as the basket payoff f smoothed by the transition density p_T in the k -th asset direction, so ϕ_k becomes smooth in x .

Vibrato Goal: Compute $\nabla_x \phi(x)$ where $\phi_k(x)$ comes from the previous CE construction.

Step 1. (Score identity)

$$\begin{aligned}\nabla_x \phi(x) &= \int f(z) \nabla_x p_T(z \mid x) dz \\ &= \mathbb{E}[f(X_T) \nabla_x \log p_T(X_T \mid x) \mid X_{N-1} \\ &= x].\end{aligned}$$

Step 2. (Vibrato variance reduction)

$$\begin{aligned}\nabla_x \phi(x) &:= \mathbb{E}\left[(f(X_T) - \phi(x)) \nabla_x \log p_T(X_T \mid x) \mid X_{N-1} = x\right]. \Leftarrow \mathbb{E}[\nabla_x \log p_T \mid x] = 0\end{aligned}$$

- CE smooths f by integrating over $p_T(\cdot \mid x)$, removing the strike kink.
- Vibrato differentiates this smooth $\phi(x)$ via the score term. The baseline $\phi(x)$ cancels constant parts \Rightarrow low variance, stable Δ .

Likelihood Ratio Method & Malliavin

LRM idea: Move derivative inside integral:

$$\begin{aligned}\frac{dF}{dx_0} &= \int f(z) \partial_{x_0} p_T(z; x_0) dz \\ &= \mathbb{E}[f(X_T) \partial_{x_0} \log p_T(X_T; x_0)].\end{aligned}$$

- The derivative acts on the smooth transition density p_T , not directly on the payoff f .
- LRM replaces pathwise differentiation of a nonsmooth f by differentiation of the smooth $\log p_T$, so the estimator remains valid across the strike kink.
- Unbiased LRM estimator for $\Delta = \frac{dF}{dx_0}$:

$$\begin{aligned}\Delta &= \mathbb{E}[f(X_T) w(X_T; x_0) \mid X_0 = x_0], \\ w(X_T; x_0) &= \partial_{x_0} \log p_T(X_T; x_0).\end{aligned}$$

Malliavin Integration by Parts

If $F = f(X_T) \in \mathbb{D}^{1,2}$ and $U \in L^2([0, T])$,

$$\mathbb{E}[\langle DF, U \rangle_{L^2}] = \mathbb{E}[F \delta(U)].$$

with $D_s F = f'(X_T) D_s X_T$. Choose U so that

$$\begin{aligned}\langle DX_T, U \rangle_{L^2} &= \partial_{x_0} X_T =: G, \\ \Delta &= \partial_{x_0} \mathbb{E}[f(X_T)] = [f'(X_T)G], \\ H_{x_0} &= \delta(U).\end{aligned}$$

Why it handles the strike:

- IBP moves the derivative from the *nonsmooth* f' to the *smooth* noise via $\delta(U)$.
- Final estimator uses $f(X_T)$ only (no f') \Rightarrow unbiased and stable across the strike kink.

Kernel & Softplus + FD

Payoff smoothing (Kernel):

$$\begin{aligned} f_\varepsilon(x) &= (f * \rho_\varepsilon)(x) \\ &= \mathbb{E}[f(x + \varepsilon Z)] \\ &= \int_{\mathbb{R}} f(x + \varepsilon z) \rho(z) dz, \\ \rho_\varepsilon(u) &= \varepsilon^{-1} \rho(u/\varepsilon). \end{aligned}$$

Smooth gradient:

$$\nabla f_\varepsilon(x) = -\frac{1}{\varepsilon} \mathbb{E}_Z [f(x + \varepsilon Z) \nabla \log \rho(Z)].$$

- Replace f by f_ε in Monte Carlo.
- Convolution removes the discontinuity of f' at $x = K$.
- ∇f_ε exists $\forall x$, bounded near K .
- Equivalent to adding small noise εZ before differentiation.

Payoff smoothing (Softplus + FD):

$$f_\alpha(x) = \frac{1}{\alpha} \log(1 + e^{\alpha(x-K)}),$$
$$f_\alpha(x) \downarrow f(x) \text{ as } \alpha \rightarrow \infty.$$

Smooth derivatives:

With $\sigma_\alpha(u) = \frac{1}{1 + e^{-\alpha u}}$, $f'_\alpha(x) = \sigma_\alpha(x - K)$,
a smooth transition from 0 to 1 near $x = K$ (instead of
the jump of $f'(x) = \mathbf{1}_{\{x>K\}}$).

- The kink in $(x - K)^+$ is replaced by a smooth region of width $\sim 1/\alpha$.
- f_α has a well-behaved curvature near the strike, so FD Greeks are stable and low-variance.
- α trades bias for variance; with finite α , small h , and enough paths, the FD Hessian converges to the Hessian of the smoothed payoff f_α .

Accuracy & Timing vs Assets

Simulation setup: $N_{\text{paths}} = 100\,000$, $N_{\text{steps}} = 50$, $\alpha = 25$.

CE / Vibrato	$n = 3$	$n = 6$	$n = 10$	$n = 15$
CE price error $ P_{\text{CE}} - P_{\text{FD}} $	1.09×10^{-2}	1.43×10^{-4}	1.18×10^{-4}	6.0×10^{-6}
Vibrato price error $ P_{\text{Vib}} - P_{\text{FD}} $	1.10×10^{-2}	1.61×10^{-4}	1.16×10^{-4}	7.0×10^{-6}
CE Hessian error $\ H_{\text{CE}} - H_{\text{FD}}\ _F$	4.38×10^{-4}	8.27×10^{-5}	1.46×10^{-4}	9.55×10^{-5}
Vibrato Hessian error $\ H_{\text{Vib}} - H_{\text{FD}}\ _F$	3.88×10^{-4}	7.31×10^{-5}	1.21×10^{-4}	8.6×10^{-5}
CE/FD time ratio	$20.45 \times$	$6.94 \times$	$2.75 \times$	$1.63 \times$
Vibrato/FD time ratio	$10.45 \times$	$3.43 \times$	$1.51 \times$	$0.84 \times$
FD calls	19	73	201	451

LRM / Malliavin / Kernel	$n = 3$	$n = 6$	$n = 10$	$n = 15$
LRM price error $ P_{\text{LRM}} - P_{\text{FD}} $	3.2×10^{-5}	5.0×10^{-5}	7.6×10^{-5}	8.5×10^{-5}
Malliavin price error $ P_{\text{Mal}} - P_{\text{FD}} $	3.2×10^{-5}	5.0×10^{-5}	7.6×10^{-5}	8.5×10^{-5}
Kernel price error $ P_{\text{Ker}} - P_{\text{FD}} $	1.6×10^{-5}	2.6×10^{-5}	4.1×10^{-5}	4.4×10^{-5}
LRM Hessian error $\ H_{\text{LRM}} - H_{\text{FD}}\ _F$	2.70×10^{-3}	7.79×10^{-1}	4.53×10^{-2}	1.293×10^0
Malliavin Hessian error $\ H_{\text{Mal}} - H_{\text{FD}}\ _F$	2.70×10^{-3}	7.79×10^{-1}	4.53×10^{-2}	1.293×10^0
Kernel Hessian error $\ H_{\text{Ker}} - H_{\text{FD}}\ _F$	9.80×10^{-5}	4.22×10^{-5}	6.64×10^{-5}	4.03×10^{-5}
LRM/FD time ratio	$5.33 \times$	$5.32 \times$	$5.23 \times$	$5.63 \times$
Malliavin/FD time ratio	$5.00 \times$	$4.94 \times$	$4.87 \times$	$5.25 \times$
Kernel/FD time ratio	$3.49 \times$	$1.52 \times$	$0.81 \times$	$0.52 \times$
FD calls	19	73	201	451

American Option Pricing Framework (MC + LSM + AAD)

1. American Option as Optimal Stopping

The value of an American option can be written as an optimal stopping problem:

$$V(t, S_t) = \sup_{\tau \in \Theta_{t,T}} \mathbb{E} \left[e^{-r(\tau-t)} \psi(S_\tau) \right],$$

where:

- $\Theta_{t,T}$ is the set of admissible stopping times,
- $\psi(\cdot)$ is the discounted payoff at exercise,
- The optimal exercise time is

$$\tau^* = \inf \{u \geq t : V(u, S_u) = \psi(S_u)\}.$$

This formulation already shows that the value depends on an *optimal boundary* separating the exercise and continuation regions.

2. Free Boundary and PDE View

In the continuous-time PDE setting, V solves an obstacle problem:

$$\max (\psi - u, u_t + \mathcal{L}u) = 0,$$

where the free boundary is the moving interface between $\{u = \psi\}$ (exercise region) and $\{u > \psi\}$ (continuation).

- At this interface the $\max(\cdot)$ creates a *kink*.
- The gradient and Hessian of u can change abruptly across the free boundary.
- This is exactly what makes American options more delicate for sensitivity analysis than European options.

Our goal is to design an AAD framework that respects this free-boundary structure, while still allowing stable second-order Greeks.

3. Longstaff–Schwartz Monte Carlo

In practice we approximate the optimal stopping rule using LSM:

$$V_t = \max \left(\text{payoff}_t, \mathbb{E}[V_{t+\Delta t} | S_t] \right),$$

where the conditional expectation is estimated via regression, e.g.

$$\mathbb{E}[V_{t+\Delta t} | S_t] \approx \beta_0 + \beta_1 S_t + \beta_2 S_t^2.$$

- Each Monte Carlo path may exercise at a different date.
- Once exercised, the path stops contributing future cashflows.
- The collection of exercise decisions approximates the free boundary in the PDE picture.

4. Why This is Hard for AD

In a naive implementation:

- The exercise rule is re-evaluated every time we bump (S_0, σ, r) .
- Tiny perturbations can change which side of the $\max(\cdot)$ a path falls on.
- The computational graph branches differently under different bumps.

Consequences for differentiation:

- The price map $V(S_0, \sigma, r)$ becomes piecewise smooth with non-differentiable kinks.
- Direct reverse-mode AD through such branching produces noisy, unstable, or even meaningless second-order Greeks.

This motivates a modification of the algorithm rather than the AD tool: the **Frozen- τ** method.

Frozen- τ Method and Smoothness

5. Frozen- τ Method (How)

We separate the problem into two stages:

- **Stage 1:** run a standard LSM algorithm once and, for each path i , record the first exercise time

$$\tau_i = \text{first } t \text{ where payoff}_t \geq \text{continuation}_t.$$

- **Stage 2:** during AAD, we reuse the same simulated paths and keep all τ_i fixed.

With τ_i frozen, the pricing formula becomes

$$V(S_0, \sigma, r) \approx \frac{1}{N} \sum_{i=1}^N e^{-r\tau_i} \psi(S_{\tau_i}^{(i)}),$$

which is a smooth function of (S_0, σ, r) given the fixed exercise times.

6. Smoothness Comparison

Aspect	No Freeze	Frozen- τ
Exercise logic	Decision flips	Fixed from 1st pass
Computational graph	Branching	Deterministic
Differentiability	Non-differentiable	Smooth
AD stability	Unstable / noisy	Stable 1st&2nd

Interpretation:

- Freezing τ removes the free-boundary effect from the AD tape.
- Any pathwise decision change is pre-captured in fixed τ_i .
- Pricing becomes smooth in (S_0, σ, r) .
- Enables stable first- and second-order Greeks via reverse-mode AAD.

AAD Engine and Scalability Settings

7. AAD Setup for Second-Order Greeks

Treat pricing parameters as AD variables:

$$x = (S_0, \sigma, r), \quad V = V(x),$$

and extract the partial Hessian:

$$H_{ij} = \frac{\partial^2 V}{\partial x_i \partial x_j}, \quad i, j \in \{S_0, \sigma, r\}.$$

Main second-order sensitivities:

$$\Gamma = V_{S_0 S_0}, \quad \text{Vanna} = V_{S_0 \sigma}, \quad \text{Volga} = V_{\sigma \sigma}.$$

Two AAD Hessian engines implemented:

- **Edge-Pushing (EP)** – recovers Hessian via adjacency list.
- **Taylor-backprop** – propagates local Taylor coefficients.

→ **Frozen- τ** ensures smooth tape, allowing both engines to compute all gradients and Hessian elements from *a single reverse sweep*.

8. Scalability Settings (for Large Monte Carlo)

Even with one reverse pass, Hessian computation over thousands of paths is memory-intensive. We control cost reduction using three mechanisms:

- ▷ **mini-batch:** Split paths into smaller batches so Hessian accumulation uses memory $\mathcal{O}(\text{batch})$ instead of $\mathcal{O}(N)$.
- ▷ **max_batch:** Hard upper bound on reverse-sweep size to prevent RAM spikes and avoid EP tape blow-up.
- ▷ **cap_active:** Remove inactive paths after Frozen- τ , reducing Hessian cost by 20–40% and improving EP stability.

These settings allow AAD-based Hessian computation to scale to tens of thousands of paths without sacrificing accuracy.

American Option Validation: Effect of cap_active




Figure 2: *

Runtime comparison of 2nd-order Greeks.

Hessian Error vs FD (2nd Order)

$$\text{rel_err}(H_{\text{EP}}, H_{\text{FD}}) = \frac{\max |H_{\text{EP}} - H_{\text{FD}}|}{\max |H_{\text{EP}}|}$$

cap_active	$\max H_{\text{EP}} - H_{\text{FD}} $	Rel. Err (%)
200	1.101×10^{-1}	0.0906
500	1.101×10^{-1}	0.0906

Interpretation

- Increasing cap_active from 200 to 500 causes the EP 2nd-order runtime to grow from 95s → 738s.
- Hessian accuracy stays unchanged (identical max error and relative error).
- Conclusion:** cap_active=200 already achieves FD-level accuracy while being **8× faster**.

American Option Validation: Mini-Batch Granularity




Figure 3: *

Runtime comparison under different mini-batch designs.

Hessian Error vs FD (2nd Order)

$$\text{rel_err}(H_{\text{EP}}, H_{\text{FD}}) = \frac{\max |H_{\text{EP}} - H_{\text{FD}}|}{\max |H_{\text{EP}}|}$$

Setting	$\max H_{\text{EP}} - H_{\text{FD}} $	Rel. Err (%)
Coarse (4000x10)	3.211×10^{-1}	0.251
Fine (1000x40)	2.348×10^{-2}	0.0169

Interpretation

- Total paths fixed at 40,000; only batching granularity changed.
- Coarse** batching runs much faster but shows larger EP–FD Hessian discrepancy.
- Fine** batching significantly improves Hessian accuracy.
- Conclusion:** fine batching for benchmark accuracy; coarse batching for speed.

Method Comparison: Bumping vs Edge-Pushing vs Taylor

Aspect	Finite (Bumping)	Difference	Edge-Pushing (EP)	Taylor Expansion
Complexity	$O(n^2)$ PDE Solved	$O(1)$ PDE solve + graph traversal	$O(1)$ forward + backward sweep	
Accuracy	Numerical errors from h	Exact (machine precision)	Exact (machine precision)	
Memory	Low (only stores values)	High (stores full tape + edge pairs)	Medium (stores (v_0, v_1, v_2) per node)	
Sparsity	N/A (no graph)	Exploits sparsity (sparse W matrix)	Limited (full monomials)	
Dense Hessian	Slow (n^2 solves)	Slower (dense W matrix)	Fast (natural for dense)	
Sparse Hessian	Slow (n^2 solves)	Fast (e.g., B-spline: 42×)	Moderate	
Implementation	Simple (black-box)	Complex (tape + graph algorithm)	Moderate (operator overloading)	
Best Use Case	<ul style="list-style-type: none"> Quick prototyping Low dimensions ($n < 5$) 	<ul style="list-style-type: none"> Sparse Hessian PDE with B-splines High dimensions 	<ul style="list-style-type: none"> Dense Hessian Deep computation graphs MC simulations 	

Project Achievements: AAD for Second-Order Greeks

